

# AI Image/Video Deepfake Detection - Roadmap

This roadmap outlines how to evolve a prototype mobile detector into a real-time, scroll-time product - suitable for VC review.

## 1 | Why You Need Vision Models, Not Just LLMs

- \* LLMs process text; they do not understand pixels.
- \* GAN and diffusion artefacts are visual - use CNN / ViT / CLIP encoders.
- \* Plan: train large vision or multimodal models, then distil / quantise for edge devices.

## 2 | Two-Tier Real-Time Architecture

React Native App

- \* 224x224 on-device thumbnail filter (tiny CNN, <10 ms)
- \* Only sends high-suspect media to heavy model

Local / Cloud Inference API

- \* FastAPI, loads larger ONNX model
- \* Returns JSON verdict + confidence

## 3 | Model-Building Phases

Phase 0 - Faces only

Dataset: 140k Real vs Fake Faces

Model: MobileNet-small (5 MB) - on-device filter

Phase 1 - Full images

Dataset: RealStock + Diffusion outputs

Model: EfficientNet-B0 (14 MB) - local API

Phase 2 - Video

Dataset: FaceForensics++, DFDC

Model: MesoNet-LSTM (25 MB) - local API

Phase 3 - Multimodal

Dataset: LAION AI-Fake mixed set

Model: Distilled CLIP tiny (30 MB) - GPU cloud

## 4 | Training Pipeline (Local)

datasets/

real/ - photographs

fake/ - GAN or diffusion

```
$ python train/train_model.py --arch mobilenet_v2 --epochs 5 --out models/mobilenet_s.onnx
```

Export to ONNX:

```
torch.onnx.export(model, dummy_input, "models/model.onnx", ...)
```

## 5 | Integration Checklist

- [ ] Gather 1k real + 1k fake images
- [ ] Train Phase-0 model, export ONNX
- [ ] Add thumbnail filter in React Native
- [ ] Build FastAPI endpoint /detect
- [ ] Log predictions for manual review

Later:

- [ ] Expand dataset to 50k+
- [ ] Train video model
- [ ] Distil INT8, benchmark 60 fps scroll
- [ ] Migrate API to GPU server

## 6 | On-Device Pre-Filter - Detailed Explanation

Purpose

- \* Acts as a bouncer: instant 224x224 scan per scroll item.
- \* Keeps UX smooth; saves battery and bandwidth.

Workflow

1. Resize frame to 224x224.
2. Tiny CNN outputs fake-probability  $p$ .
  - $p < 0.3$  -> mark safe.
  - $p > 0.7$  -> send full-res to heavy model.
3. Heavy model corrects any false alarms.

Model Choices

- \* MobileNet-V2 tiny (3 MB, 5-8 ms)
- \* ShuffleNet-V2 0.5x (2.5 MB, 4-6 ms)
- \* MesoNet-lite (1 MB, ~3 ms)

Training Tips

- \* Use same dataset but resized and heavily augmented.
- \* Stop at ~85% accuracy; heavy model is safety net.
- \* Quantise to INT8 for TFLite / ONNX.

Threshold Tuning

- \* Safe cut-off: false positive rate  $\leq 3\%$ .
- \* Suspect cut-off: true positive rate  $\geq 85\%$ .

Performance Proof (Pixel 5)

- \* 7 ms per frame, 120 thumbnails / s, 30% CPU used.

Limitations & Mitigations

- \* Might miss subtle fakes - send uncertain frames to heavy model.
- \* Thumbnail blur - sharpen lightly before inference.

## **7 | Where an LLM Fits Later**

- \* Generate user-friendly explanations of detected artefacts.
- \* Cluster false positives for rapid retraining.
- \* Cross-modal sanity checks: compare Whisper transcript with lip movement.

## **End of Roadmap**

Questions? Contact: [your-email@company.com](mailto:your-email@company.com)